ONLY YOU — one-page brief

The more capable the model, the more it needs an authority layer.

ONLY YOU is the authority layer that lets a wearable agent act in the world without being trusted with it. Sushant Poudel · Meta Global AI Developer Hackathon 2026 · Build for Your World

The problem

Meta's WebMCP documentation says the untrusted-content hint is *"not an enforced security boundary,"* and the assistant cannot tell the wearer's voice from a stranger's. With a camera, microphone and wallet, any TV ad, poster, webpage or invoice can issue commands. Capability is solved; **authority is not.**

What it does

Provenance	Only the enrolled owner's voice is a command. Scene text, page text and documents are data — enforced by channel, not by prompt wording.
Bounded delegation	Ed25519-signed mandates with caps, per-merchant/daily sub-limits, recurring allowances and sub-agents that can only narrow authority. A merchant verifies one offline with only the public key — or a QR code.
Provable evidence	Every action *and every refusal* is a signed, hash-chained receipt. Export the audit bundle and re-verify it in Python or **in the browser**; every receipt has a public verification page.

The evidence

Measurement	Result
Confused-deputy corpus, Muse Spark 1.3	undefended **25%** executed · prompt-only 37.5% · provenance-only 50% · **full gate 0%**, 0% false refusals
Same corpus, contributor build	50% → 37.5% → 50% → **0%**
Post-build reproduction on final code	every arm **0%**, **16/16 receipts**, chain intact, 99 live calls
Label layer (the part that can be wrong)	DOM recall 81.0%, 14.3% false positives — published, with 8 missed techniques
Forged sensor signals	6/6 executed without a signed bus → **0/6** with it
Cross-model audit	same gate + receipt format under Muse Spark and a local model, verified in the browser
Engineering	**172 tests**, 67-check self-audit, CI (tests + Docker), four committed real-model runs

Why it matters commercially

- **Wearables**: the wearer-vs-the-world trust story is the product question for glasses.
- **Agentic commerce**: merchants need to verify an agent's authority before charging; Amazon blocks Muse today.
- **Enterprise copilots**: the same gate gives delegated, receipted, revocable action instead of blind trust.

Try it in 30 seconds

- Playground (no install): <https://sushant-me.github.io/only-you/playground.html>
- Verify an audit bundle: <https://sushant-me.github.io/only-you/verify.html>
- Repo / one command: `make judge` · `docker run -p 8000:8000 only-you-demo`
- 3-minute narrated demo, report and deck: <https://sushant-me.github.io/only-you/>

The honest line

The gate is engineered and tested; the label layer is measured and published where it fails. Cloned voice, an exposed demo server, cross-device policy trust and synthetic narration are named in `limitations.md` — not hidden. That is what makes this auditable rather than marketing.

Ask: a pilot with the wearables or agentic-commerce team — and a role building this class of infrastructure at Meta.