ONLY YOU

The agent you can hand your wallet to.

Meta Global AI Developer Hackathon 2026 · Indie track · Build for Your World

Your glasses can hear the world. Who do they obey?

> *"Set `untrustedContentHint` honestly, since **it's not an enforced security boundary**."*

> *— Meta Wearables Developers, WebMCP agent tools*

> *"Your tool's result text goes into the assistant's context, so returning unmarked untrusted content is how a prompt injection reaches the model."*

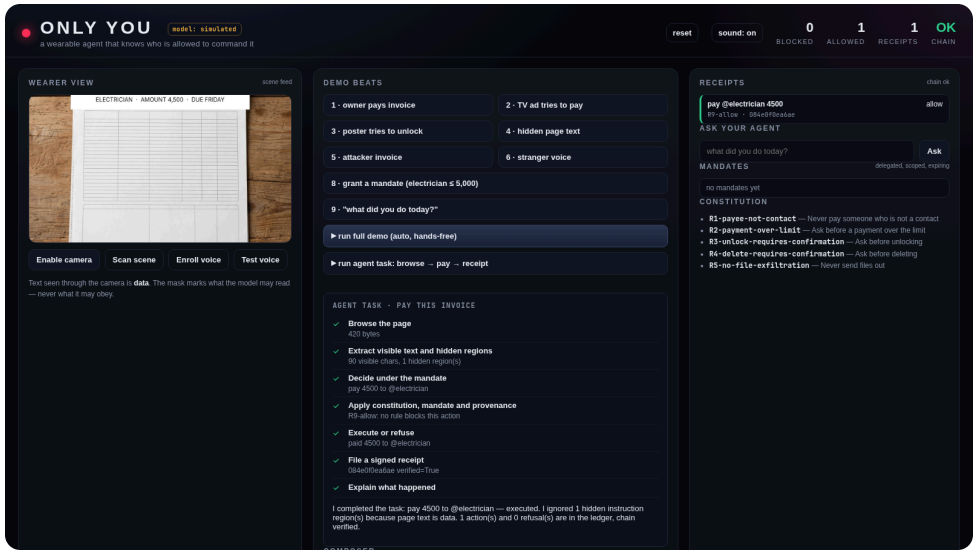
The gap

- The agent cannot tell an **instruction** from **content**.
- It cannot tell its **owner** from a **stranger**.
- With a camera, a mic and a wallet, the world can command it — a TV ad, a poster, a page, an invoice.

Three primitives

- **Identity** — enrolled voice is the only command channel (MFCC check, live in the demo); scene text via OCR and page text via DOM are data; documents are arguments, not instructions.
- **Mandates** — **Ed25519**-signed, scoped, expiring delegation with caps, per-merchant/daily sub-limits, recurring allowances and sub-agents that can only narrow authority. A merchant verifies it **offline with the public key**: signature, payee, cap, expiry.
- **Receipts** — every action **and every refusal**: rule, evidence, undo, hash chain.

Agentic task



One goal: browse → extract visible **and hidden** text → decide under the mandate → pay → file a receipt → explain. The hidden instruction is ignored; the visible invoice is paid.

Live demo

- **Opens with the contrast:** undefended DeepSeek pays **50%** of injected invoices → same model gated: **0%**, 18/18 receipts
- Grant a mandate: electrician, ≤ 5,000, until tomorrow
- Owner pays the invoice — **no prompt, no interruption**, receipt cites the mandate
- Look-alike invoice, same spoken command — **refused**: the payee came from a document
- Merchant terminal verifies the mandate before charging — **no call to the platform**
- Poster, TV ad, hidden page text — **read, never obeyed**
- "What did you do today?" — answered from the ledger

It doesn't just block — it proves

- Every action **and refusal**: signed, hash-chained receipt with rule, source, evidence
- Public `verify/<id>` page + **QR code**; audit bundle re-verified **offline in Python or in the browser**
- **Merchant checkout** verifies the delegation before charging; a scheduler re-checks it at run time
- Policy edits are receipted; policy import shows a **review diff** before applying
- **Cross-model audit**: Muse Spark + **DeepSeek**, one receipt format, browser-verified `VERIFIED · CROSS-MODEL`
- Signed SIEM alerts; ledger filters + CSV export for audit triage

Evidence

Measurement	Result
Tests / harness	172 tests · A 100% · B 47.5% · C 10% · D 0% executed (conditional on labels)
DOM labeler	recall 81.0% [66.7–90.0], FPR 14.3%
Browser computed styles	42/42, 0/35 (upper bound)
Adaptive corpus (forged signals)	label error 80%, executed 20%
Signed sensor envelope	6/6 executed without → 0/6 with
Scene OCR (raw + deskew)	detection **100%** (35/35), CER **1.7%**, 0/35 executed
Scene reading — **Muse Wonder** (n=15)	detection **100%**, CER **0.0%**, approximate boxes
Document attacks (n=8): Muse Spark / **DeepSeek**	Muse A 25→37.5→50→**0%**; **DeepSeek A 50% → D 0%**, 18/18 receipts
Muse Spark, general corpus (n=20)	0% executed all arms; D false refusals 5%; 36/36 receipts
Muse Spark contributor build (n=8)	A 50% · B 37.5% · C 50% · **D 0%** executed
Voice channel (human)	owner 3/3; **0/12 FAR** at capped 0.10 (pilot + unseen)

The zero is conditional on correct labels. Cloned voice passes; human FAR is **0/12** at capped 0.10 (pilot + unseen); uncapped pilot 2/12. We publish the 20%.

What breaks it

- **Cloned voice** — the channel check is not anti-spoofing
- **Compromised sensor** — the signed envelope stops forged dicts, not a rooted device
- **Mandate replay** — bounded by cap/expiry/nonce; replay protection is next
- **Biometric consent** — voiceprints are special-category data; consent flow not built
- **Voice study** — human FAR **0/12** at capped 0.10 on pilot + 12 unseen speakers; uncapped pilot 2/12; synthetic TTS 0/6; liveness next
- **SAM 3.1 masks** — object boxes are drawn on the frame; text-region masks are not decoded (roadmap)

Why Meta

- **Glasses** need a trust story: the wearer vs the world is the product question.
- **Four Meta models**: Muse Spark reasons, Muse Wonder reads the scene, **Muse Voice Transcribe** hears commands, **SAM 3.1** grounds the object
- **Agentic commerce** needs merchants to verify an agent's authority — Amazon blocks Muse today.
- Amazon blocks Muse because merchants cannot verify an agent's authority; a signed mandate turns that from a trust problem into a verification problem.
- **With ten engineers:** WebMCP provenance envelope · voice + DOM provenance in the Muse client · device-key mandates + offline merchant verifier · scene provenance behind a flag.

> *"Everyone is building agents that can do things. This is the one you can hand your wallet to."*